

Laboratory 2 Report

ANDREA BOTTICELLA*, s347291, Politecnico di Torino, Italy

ELIA INNOCENTI*, s345388, Politecnico di Torino, Italy

SIMONE ROMANO*, s344024, Politecnico di Torino, Italy

This work addresses malware detection through behavioral analysis of API call sequences. We implement and compare four classification approaches: a frequency-based logistic regression baseline, Feed-Forward Neural Networks with learned embeddings, Recurrent Neural Networks (RNN, BiRNN, LSTM) for sequential modeling, and Graph Neural Networks (GCN, GraphSAGE, GAT) that represent call sequences as transition graphs. Experiments are conducted on a dataset of execution traces containing 258 unique API calls, with attention to class imbalance handling and variable-length sequence processing.

1 INTRODUCTION

Dynamic malware analysis leverages runtime behavior to detect malicious software, circumventing the limitations of static signature-based methods against obfuscated or polymorphic variants. API call sequences—the ordered invocations a program makes to the operating system—provide a rich behavioral fingerprint, capturing memory operations, registry accesses, and system queries that reveal malicious intent.

This laboratory implements a comparative study of neural network architectures for binary malware classification based on API call sequences. The dataset consists of execution traces, each represented as a variable-length sequence of API function names (e.g., `NtAllocateVirtualMemory`, `RegOpenKeyExW`).

We evaluate four progressively complex approaches:

- **1:** Implement and evaluate a frequency-based logistic regression baseline.
- **2:** Develop Feed-Forward Neural Networks with different input representations.
- **3:** Explore Recurrent Neural Network architectures for sequential modeling.
- **4:** Construct Graph Neural Networks to leverage structural information in API call sequences.

The following sections detail the methodology and results for each task, analyzing how architectural choices impact classification performance.

2 TASK 1: FREQUENCY-BASED BASELINE

This section establishes a baseline for malware classification using a frequency-based representation of API call sequences. The intent is to quantify how far one can go with a simple model that ignores temporal order, and to provide a reference point for the more expressive architectures introduced later.

2.1 Data Loading and Exploration

The dataset is provided in JSON format, with each sample containing two fields: `api_call_sequence` (a list of API function names observed during execution) and `is_malware` (a binary label). Both training and test partitions are loaded for an initial inspection of size, label distribution, and vocabulary.

The training set contains 16,325 samples, while the test set comprises 6,505 samples. A striking observation emerges from examining the class distribution: malware samples constitute 96.26% of both partitions (15,714 training and 6,262 test samples), leaving only 3.74% for benign software (611 training and 243 test samples). This severe class imbalance has important implications for model evaluation. Accuracy alone would be misleading, as a trivial classifier predicting “malware” for every sample would achieve 96% accuracy without learning anything meaningful. We therefore pay close attention to precision, recall, and F1-score for both classes, particularly monitoring performance on the minority benign class.

2.2 Vocabulary Extraction and Analysis

The frequency-based representation requires a fixed vocabulary. To avoid any information leakage, the vocabulary is extracted from the training set and used to define the feature space at inference time. The training set contains 258 unique API calls, while the test set contains 232.

This immediately exposes a vocabulary mismatch. Twenty-nine API calls appear only in the training set (they can influence training but never occur at test time), while three appear exclusively in the test set: `NtDeleteKey`,

*The authors collaborated closely in developing this project.

`ControlService`, and `WSASocketA`. The latter case is the critical one, because the classifier would be asked to process signals that were not available during training.

Unseen API calls can be handled by introducing an explicit unknown feature (e.g., `<UNK>`) or by discarding them when building the input vector. Since the three unseen calls appear in only three test samples, they have negligible impact at the dataset scale; the pipeline therefore ignores them and builds test vectors strictly over the training vocabulary.

2.3 Frequency-Based Feature Vectors

Each API call sequence is transformed into a 258-dimensional vector, where each element is the count of occurrences of a specific API call from the training vocabulary. This bag-of-words representation discards temporal ordering: any permutation of the same multiset of calls maps to the same vector, and patterns that depend on call order are not represented.

The resulting feature matrices are highly sparse, as individual execution traces invoke only a small subset of the full API vocabulary. On average, training samples contain 21.95 non-zero elements per row, representing just 8.51% of the 258 dimensions. Test samples show slightly higher density with 24.28 non-zero elements (10.46%). This sparsity has practical implications: memory-efficient sparse matrix representations can be leveraged, and classifiers designed for high-dimensional sparse data become particularly suitable choices.

2.4 Classifier Selection and Training

Logistic regression is adopted as the baseline classifier for its computational efficiency on sparse high-dimensional inputs and its transparency: the learned coefficients provide a direct indication of which API calls are associated with either class. We use the `liblinear` solver, which is well suited to small and sparse binary classification problems, with the default L2 regularization strength ($C=1.0$), no class weighting, and `random_state=42` for reproducibility. We deliberately avoid extensive hyperparameter tuning at this stage: the goal is to obtain a simple, reproducible reference point rather than optimize the baseline aggressively.

A baseline is not a formality. It anchors the comparison with the more complex architectures in the following tasks and prevents attributing improvements to architectural sophistication when the same gains could be achieved by a simpler representation or by exploiting label imbalance.


The data is split into 60% training and 40% validation using stratified sampling to preserve class proportions. Training produces a log-loss of 0.0806 on the training set and 0.1225 on the validation set. The gap between these values suggests mild overfitting, but remains acceptable for a baseline model using only default L2 regularization and no task-specific tuning.

2.5 Evaluation and Results

We evaluate the trained classifier on both validation and held-out test sets. Table 1 presents the full classification report for the test set.

Table 1. Frequency-based baseline: classification report (test set)

Class	Precision	Recall	F1-Score	Support
Benign (0)	0.6885	0.3457	0.4603	243
Malware (1)	0.9751	0.9939	0.9844	6,262
Accuracy		0.9697		6,505
Macro avg	0.8318	0.6698	0.7223	6,505
Weighted avg	0.9644	0.9697	0.9648	6,505



(a) Confusion matrix for the frequency-based baseline on the test set.

(b) Confusion matrix for the frequency-based baseline on the validation set.

Fig. 1. Comparison of confusion matrices for the frequency-based baseline on test and validation sets.

The results reveal a markedly asymmetric behavior driven by the underlying imbalance. Malware detection is extremely strong (precision 97.51%, recall 99.39%, F1-score 0.9844), and the confusion matrices in Figure 1 confirm that false negatives are rare. The benign class is more challenging: although benign predictions are often correct (precision 68.85%), recall is low (34.57%), meaning that a large fraction of benign samples are flagged as malware. With only 3.74% benign samples, the model learns a conservative decision boundary that favors catching malware at the cost of additional false alarms.

Despite ignoring order and operating on sparse vectors, the frequency-based baseline reaches 96.97% accuracy and near-perfect malware recall. This shows that the presence and relative frequency of specific API calls already carries substantial discriminative signal, even before exploiting sequential dependencies. At the same time, the macro F1-score of 0.7223 and the benign recall of 34.57% make clear that the model is not equally reliable across classes: its high accuracy is mainly driven by the dominant malware class.

The limitation is structural: frequency vectors cannot represent temporal motifs, repeated call patterns, or long-range dependencies. This motivates the sequence- and structure-aware models in the subsequent tasks, where the representation is designed to preserve more of the behavioral dynamics that are erased by the bag-of-words transformation.

3 TASK 2: FEED FORWARD NEURAL NETWORK (FFNN)

The Feed-Forward Neural Network (FFNN) represents the first deep learning architecture applied to the detection task, moving beyond bag-of-words counts to a positional representation of API calls. This transition introduces two primary constraints: the requirement for fixed-size input vectors and the lack of explicit temporal modeling.

3.1 Sequence Length Variability and Fixed-Size Encoding

The API call traces exhibit significant variability in length. An analysis of the training set shows a distribution ranging from 60 to 90 calls (mean 75.03), while the test set is shifted toward higher values, ranging from 70 to 100 calls (mean 86.33). This mismatch is visualized in Figure 2. Therefore, the test set contains sequence lengths that are never observed during training, which is particularly relevant for a fixed-input architecture.




Fig. 2. Distribution of API call sequence lengths in training and test sets.

To accommodate the fixed input dimension of the FFNN, we established a target length of **90** based on the maximum observed in the training data. This choice ensures that no training information is lost via truncation. Sequences shorter than 90 are extended using **zero-padding**, while test sequences exceeding this limit are **truncated**.

The fixed length is estimated only from the training partition, in order to avoid using information from the held-out test set to shape the preprocessing pipeline. Alternative choices such as the median length (75) or the 75th percentile (83) would reduce memory usage, but would also discard part of the training sequences. Since the maximum training length is not an extreme outlier, we choose the conservative value of 90.

3.2 Categorical Encoding

Before feeding sequences to the network, API names must be mapped into a numerical representation. The vocabulary is extracted from the padded training sequences, assigning index 0 to `<PAD>` and index 1 to `<UNK>`; the resulting vocabulary contains 260 entries, including the two special tokens. Unseen API calls in the test set are mapped to `<UNK>`.

We evaluate two alternatives. In the **sequential ID** setting, each API call is represented directly by its integer identifier, producing a 90-dimensional input vector. This is simple but problematic: the model can interpret IDs as ordered numerical quantities, even though the numbering is arbitrary. In the **learnable embedding** setting, each API ID is first mapped to a 128-dimensional trainable vector, and the 90 embedded calls are flattened before entering the dense layers. This gives the model the opportunity to learn a more meaningful representation of API calls.

3.3 Architectural Design and Hyperparameters

We implemented two variants of the FFNN to evaluate the impact of input representation:

- (1) **Sequential IDs:** Raw integer identifiers are fed directly into a deep network with hidden dimensions of [256, 128, 64]. This model treats IDs as ordinal values, which is an architectural limitation.
- (2) **Learnable Embeddings:** Each API ID is mapped to a **128-dimensional** continuous vector space. This allows the model to learn semantic similarities between calls. The hidden layers were set to [128, 64, 32].

Both models utilize **Batch Normalization** for training stability and **Dropout** (0.3–0.4) to mitigate overfitting. Training is performed with **AdamW**, **BCEWithLogitsLoss**, batch size 64, and early stopping with patience 15 on validation loss. Since label 1 corresponds to malware and is the majority class, the loss uses `pos_weight=0.20` to down-weight malware errors and prevent the minority benign class from being completely overwhelmed.

For the sequential-ID model, we test learning rates $\{10^{-4}, 5 \cdot 10^{-5}, 10^{-5}\}$ with weight decay 10^{-5} . For the embedding model, we test $\{5 \cdot 10^{-5}, 10^{-5}, 5 \cdot 10^{-6}\}$ with weight decay 10^{-4} . The final configuration is selected by the minimum validation loss, while the macro F1 curves are used as a consistency check for balanced performance under class imbalance.




Fig. 3. Validation loss comparison across learning rates for the two FFNN input representations.




Fig. 4. Validation macro F1 comparison across learning rates.

3.4 Results and Discussion

Table 2 summarizes the selected models. The embedding-based FFNN obtains a lower validation loss and better validation macro F1 than the sequential-ID model, consistently with Figure 4. The sequential-ID curves are less informative because the network can exploit accidental numerical magnitudes of API identifiers; embeddings reduce this artifact by learning continuous API-call representations. On the test set, embeddings improve benign recall from 11.93% to 30.86%, but the minority class remains difficult because the model still receives a fixed, truncated vector rather than a temporal trace.

Table 2. Task 2 FFNN results for the selected configurations.

Model	LR	Min Val Loss	Val Macro F1	Test Acc.	Test Macro F1
Sequential IDs	10^{-4}	0.0919	0.5624	0.9301	0.5383
Learnable Embeddings	$5 \cdot 10^{-5}$	0.0726	0.6662	0.9184	0.5886

Overall, both FFNN variants remain below the frequency-based baseline of Task 1 in macro F1. The main reason is architectural: after padding or truncation, the FFNN cannot model the temporal evolution of the API trace, and 40% of test sequences are longer than 90 calls, so useful late-sequence behavior may be removed. This motivates the recurrent models in the next task.

4 TASK 3: RECURRENT NEURAL NETWORK (RNN)

This section investigates recurrent neural networks for malware classification from ordered API call traces. Differently from the FFNNs in Task 2, recurrent models process the sequence step by step and can exploit temporal dependencies without forcing all samples into the same global length.

4.1 Sequence Preprocessing

RNNs still require padding to build mini-batches, but padding is handled dynamically rather than globally. We use a custom `collate_fn` that sorts the samples in each batch by decreasing length, pads them only up to the longest sequence in that batch, and returns both the padded tensor and the original lengths. The embedding layer reserves

index 0 for padding through `padding_idx=0`, and `pack_padded_sequence` is used before the recurrent layer. In this way, the model does not update its hidden state on dummy positions and avoids wasting computation on padding tokens.

This preprocessing also removes the need for test-time truncation. A recurrent network shares the same parameters across time steps, therefore longer traces simply imply additional recurrent updates. This is an important difference with respect to Task 2, where test sequences longer than 90 API calls had to be truncated to match the FFNN input size. As a consequence, recurrent models can preserve late API calls that may contain useful behavioral information.

From a memory perspective, dynamic padding makes the batch tensor proportional to the longest sequence in the current batch, not to a fixed global maximum. Moreover, recurrent models keep a compact 32-dimensional embedding for each API call and process it sequentially, instead of flattening the whole embedded sequence into a large dense vector as in the embedding-based FFNN. The price is computational speed: the recurrent dependency between time steps reduces parallelization and makes training slower than FFNN training.

4.2 RNN Architectures and Model Selection

We trained the three recurrent variants required by the assignment: a simple one-directional RNN, a bidirectional RNN, and a bidirectional LSTM. All models use a vocabulary built from the training set, 32-dimensional embeddings, batch size 64, AdamW optimization, binary cross-entropy with logits, early stopping, and a plateau-based learning-rate scheduler. Since label 1 corresponds to malware and is the majority class, `pos_weight=0.20` down-weights malware errors and helps the minority benign class influence the optimization.

The Simple RNN uses two unidirectional recurrent layers with hidden size 128, `tanh` activation, and dropout $p = 0.3$. The bidirectional RNN keeps the same hidden size but processes each sequence in both directions, using dropout $p = 0.35$. Finally, the bidirectional LSTM replaces the vanilla recurrent units with gated memory cells and uses dropout $p = 0.4$, with the goal of improving stability on longer dependencies. For the three architectures we tested, respectively, learning rates $\{2 \cdot 10^{-4}, 10^{-4}, 8 \cdot 10^{-5}, 5 \cdot 10^{-5}\}$, $\{3 \cdot 10^{-4}, 10^{-4}, 8 \cdot 10^{-5}, 5 \cdot 10^{-5}\}$, and $\{5 \cdot 10^{-4}, 3 \cdot 10^{-4}, 10^{-4}, 8 \cdot 10^{-5}\}$.




Fig. 5. Validation curves for the Simple RNN across learning rates.




Fig. 6. Validation curves for the bidirectional RNN across learning rates.




Fig. 7. Validation curves for the bidirectional LSTM across learning rates.

The final configurations are selected by comparing validation loss and validation macro F1, as in Task 2. The Simple RNN performs best with $lr = 8 \cdot 10^{-5}$, reaching a minimum validation loss of 0.0423 and validation macro F1 of 0.7824. The bidirectional RNN also selects $lr = 8 \cdot 10^{-5}$, with a slightly lower validation loss of 0.0406 and macro F1 of 0.7861. The bidirectional LSTM is the strongest validation model: with $lr = 3 \cdot 10^{-4}$ it reaches the lowest validation loss, 0.0288, and the highest validation macro F1, 0.8346.

RNNs are significantly slower than the FFNNs from Task 2. The selected FFNN configurations train in roughly 38–39 seconds, while the selected Simple RNN requires about 4 minutes and 45 seconds, the bidirectional RNN about 8 minutes and 42 seconds, and the bidirectional LSTM about 5 minutes and 3 seconds. This overhead is caused by recurrent step-by-step processing, backpropagation through time, dynamic sorting and padding inside the data loader, and packed-sequence handling.

4.3 Performance Analysis

Table 3 reports the test classification metrics of the selected recurrent models and compares them with the previous baselines. The Simple RNN already improves the minority-class recall substantially, increasing benign recall from 30.86% for the best FFNN to 70.78%. This confirms that the order of API calls contains useful information that is mostly lost by fixed-window FFNNs.

Table 3. Performance comparison on the test set.

Model	Accuracy	Benign (0)		Malware (1)		Macro F1
		Prec.	Rec.	Prec.	Rec.	
Simple RNN	95.34%	0.4257	0.7078	0.9884	0.9630	0.7536
Bi-RNN	95.56%	0.4392	0.6831	0.9874	0.9661	0.7556
Bi-LSTM	96.66%	0.5387	0.7449	0.9899	0.9752	0.8039
Freq. baseline	96.97%	0.6885	0.3457	0.9751	0.9939	0.7223
Seq. ID FFNN	93.01%	0.1074	0.1193	0.9657	0.9615	0.5383
Embed. FFNN	91.84%	0.1712	0.3086	0.9723	0.9420	0.5886

The bidirectional RNN only marginally improves macro F1 with respect to the Simple RNN, despite requiring more training time. Its backward pass provides additional context, but the gain is limited on this dataset. The bidirectional LSTM is instead clearly better: its gates stabilize the recurrent state and help preserve relevant information over longer API traces, producing the highest benign recall and the best macro F1 among the Task 3 models.

The validation curves in Figures 5–7 show that recurrent training is less smooth than FFNN training. Small changes in learning rate produce visible changes in convergence, and vanilla RNNs are more exposed to vanishing or unstable gradients. The LSTM is the most stable architecture among the tested variants, because its input, forget, and output gates regulate how information is stored and discarded across time.

The frequency baseline still obtains the highest accuracy, because it is very strong on the majority malware class. However, its benign recall is only 34.57%, meaning that many benign samples are classified as malware. The bidirectional LSTM lowers accuracy slightly, from 96.97% to 96.66%, but improves macro F1 from 0.7223 to 0.8039 and benign recall from 34.57% to 74.49%. In an imbalanced setting, this is a better balanced behavior: the model sacrifices a small amount of majority-class accuracy while substantially reducing the asymmetry between the two classes.

5 TASK 4: GRAPH NEURAL NETWORK (GNN)

The last modeling strategy represents each API call sequence as a graph. While RNNs preserve the execution trace as a temporal chain, GNNs reinterpret the same behavior as a set of API calls connected by transitions. This allows the model to reason on local neighborhoods and repeated structural patterns, rather than only on a fixed vector or on a left-to-right hidden state.

5.1 Graph Construction and Preprocessing

Each sample is converted into a transition graph. The nodes correspond to the unique API calls appearing in the sequence, while edges connect consecutive calls with window size 1. Therefore, a transition such as $A \rightarrow B$ creates an edge from the local node representing A to the local node representing B ; loops are naturally allowed when the same API call follows itself. Node features are the global API identifiers, which are then mapped to 32-dimensional learnable embeddings inside the GNN. Unknown API calls are mapped to the `<UNK>` index, and padding is excluded from the graph construction.

This representation removes the need for padding. Unlike FFNNs and RNNs, PyTorch Geometric batches graphs by forming a larger disconnected graph and keeping a `batch` vector that records the graph membership of each node. As a result, each sample can keep its own number of nodes and edges, and the model produces a graph-level representation through `global_mean_pool`. In the first training graph, for example, only 21 unique API-call nodes and 80 transition edges are needed, even though the global vocabulary contains 260 entries.


The same argument explains why test sequences are not truncated. Longer traces simply generate additional transitions and possibly additional nodes. Truncating them would remove edges from the behavioral graph and would partially defeat the reason for using a graph representation. This is a clear advantage over the FFNN setting of Task 2, where sequences longer than 90 calls had to be cut.

Compared with RNNs, GNNs can aggregate information from neighborhoods and capture recurring transition motifs, such as an API call connected to several different successors or repeated cycles between calls. This can make the representation more robust to small local reorderings. However, this comes at a cost: after graph construction and global pooling, the model no longer knows exactly whether a transition occurred at the beginning or at the end of the trace. Thus, GNNs preserve structural relations between API calls, but lose part of the strict chronological information explicitly modeled by RNN hidden states.


5.2 GNN Architectures and Model Selection

We trained three graph architectures: GCN, GraphSAGE, and GAT. All models share the same basic setup: 32-dimensional embeddings, hidden dimension 64, three graph convolution/message-passing layers, global mean pooling, batch size 32, AdamW optimization, early stopping, and binary cross-entropy with logits. As in the previous neural models, the loss uses `pos_weight=0.20`, since malware is the positive and majority class; this down-weights malware errors and gives more relative influence to the minority benign samples.

The GCN model uses three `GCNConv` layers followed by ReLU activations and dropout $p = 0.2$, then a small classifier with dropout $p = 0.5$. GraphSAGE replaces the convolutional operator with three `SAGEConv` layers, allowing node representations to be updated through neighborhood aggregation. GAT uses three attention-based layers: the first `GATConv` uses 4 attention heads, while the following layers use a single head before the graph-level pooling. For GCN and GraphSAGE we tested learning rates $\{10^{-4}, 5 \cdot 10^{-4}, 10^{-3}\}$, while for GAT we tested $\{5 \cdot 10^{-5}, 10^{-4}, 5 \cdot 10^{-4}\}$.




(a) GCN validation loss.




(b) GCN macro F1.

Fig. 8. Validation curves for GCN across learning rates.




(a) GraphSAGE validation loss.




(b) GraphSAGE macro F1.

Fig. 9. Validation curves for GraphSAGE across learning rates.



(a) GAT validation loss.



(b) GAT macro F1.

Fig. 10. Validation curves for GAT across learning rates.




Fig. 11. Validation macro F1 comparison among the selected GNN architectures.

The final configurations are selected using validation behavior, consistently with the previous tasks. For GCN, $lr = 5 \cdot 10^{-4}$ is selected because it achieves the best validation-loss minimum among the tested configurations, with validation macro F1 0.7830 and test macro F1 0.7711. GraphSAGE also selects $lr = 5 \cdot 10^{-4}$ under the same validation-loss criterion, reaching test macro F1 0.7633. Finally, GAT selects $lr = 5 \cdot 10^{-4}$, obtaining validation macro F1 0.8079 and the best GNN test macro F1, 0.7903. The selected training times are approximately 179.62 seconds for GCN, 135.01 seconds for GraphSAGE, and 257.52 seconds for GAT.

5.3 Performance Analysis

Table 4 provides the final comparison across all selected architectures. The GNNs clearly outperform both FFNN variants from Task 2, confirming that transition structure is much more informative than fixed positional vectors. They also improve over the frequency baseline in macro F1, although the baseline remains very competitive in accuracy because it is extremely strong on the majority malware class.

Table 4. Final performance comparison on the test set.

Model	Accuracy	Benign (0)		Malware (1)		Macro F1
		Prec.	Rec.	Prec.	Rec.	
GCN	95.82%	0.4619	0.7243	0.9891	0.9673	0.7711
GraphSAGE	95.10%	0.4191	0.8107	0.9924	0.9564	0.7633
GAT	96.31%	0.5042	0.7407	0.9898	0.9717	0.7903
Simple RNN	95.34%	0.4257	0.7078	0.9884	0.9630	0.7536
Bi-RNN	95.56%	0.4392	0.6831	0.9874	0.9661	0.7556
Bi-LSTM	96.66%	0.5387	0.7449	0.9899	0.9752	0.8039
Freq. baseline	96.97%	0.6885	0.3457	0.9751	0.9939	0.7223
Seq. ID FFNN	93.01%	0.1074	0.1193	0.9657	0.9615	0.5383
Embed. FFNN	91.84%	0.1712	0.3086	0.9723	0.9420	0.5886

Among the GNN models, GAT is the strongest overall, suggesting that attention helps assign different importance to neighboring API-call transitions. GCN is close behind, while GraphSAGE reaches the highest benign recall, 81.07%, at the cost of lower benign precision. All three models keep excellent malware precision and recall, above 98.9% and 95.6% respectively.

The comparison with the baseline highlights the same trade-off already observed in Task 3. The frequency-based classifier reaches 96.97% accuracy and 99.39% malware recall, but only 34.57% benign recall. GNNs sacrifice a small amount of accuracy while raising benign recall to the 72–81% range, producing better macro F1 and a more balanced behavior.

Compared with the recurrent models, the best GNN remains slightly below the Bi-LSTM in macro F1: GAT reaches 0.7903, while the Bi-LSTM reaches 0.8039. This suggests that strict temporal information is still useful, although graph models remain competitive by summarizing transition structure and clearly outperforming fixed-length FFNNs.